

1. Code with detailed explanations (30%)

i. Part1 (kernel k-means 5%, normalized cut 5%, ratio cut 5%)

```
def load_image(image_path):
    img = np.array(Image.open(image_path).convert('RGB'))
    h, w, _ = img.shape
    return img, h, w
```

kernel k-means

$$\mathbf{k}(x, x') = e^{-\gamma_s \|S(x) - S(x')\|^2} \times e^{-\gamma_c \|C(x) - C(x')\|^2}$$

- $S(x)$: spatial coordinate of pixel x (e.g., row, column)
- $C(x)$: color (RGB) value of pixel x
- γ_s : controls the influence of spatial distance
- γ_c : controls the influence of color difference

```
def compute_kernel_matrix(img, gamma_s, gamma_c):
    h, w, c = img.shape
    N = h * w

    # Create spatial coordinate vectors: S(x)
    spatial = np.array([[i // w, i % w] for i in range(N)])

    # Flatten the image to get color vectors: C(x)
    rgb = img.reshape((-1, 3))

    # Compute pairwise squared distances in spatial domain: ||S(x) - S(x')||^2
    spatial_dists = squareform(pdist(spatial, 'sqeuclidean'))

    # Compute pairwise squared distances in color domain: ||C(x) - C(x')||^2
    rgb_dists = squareform(pdist(rgb, 'sqeuclidean'))

    # Combine both distance matrices using Gaussian kernel
    #  $k(x, x') = \exp(-\gamma_s * \text{spatial\_dist}) * \exp(-\gamma_c * \text{color\_dist})$ 
    K = np.exp(-gamma_s * spatial_dists) * np.exp(-gamma_c * rgb_dists)

    return K
```

The squared distance between a sample x_j and the centroid μ_k^ϕ of cluster k in the feature space is:

$$\|\phi(x_j) - \mu_k^\phi\|^2 = \left\| \phi(x_j) - \frac{1}{|C_k|} \sum_{n=1}^N \alpha_{kn} \phi(x_n) \right\|^2$$

Expanding this using kernel functions:

$$= \underbrace{\mathbf{k}(x_j, x_j)}_{selfsimilarity} - \frac{2}{|C_k|} \sum_n \alpha_{kn} \mathbf{k}(x_j, x_n) + \frac{1}{|C_k|^2} \sum_p \sum_q \alpha_{kp} \alpha_{kq} \mathbf{k}(x_p, x_q)$$

- $\phi(x)$: The high-dimensional mapping of x
- μ_k^ϕ : The centroid of cluster k in the feature space
- α_{kn} : 1 if $x_n \in C_k$; otherwise 0
- $\mathbf{k}(x_i, x_j)$: The kernel function

```

def kernel_kmeans(K, k, max_iter=100, tol=1e-3, init_method='kmeans++'):
    N = K.shape[0]
    if init_method == 'random':
        np.random.seed(42)
        labels = np.random.randint(0, k, size=N)
    elif init_method == 'kmeans++':
        np.random.seed(42)
        centers = [np.random.randint(0, N)]
        for _ in range(1, k):
            D = np.array([1 - np.max(K[i, centers]) for i in range(N)])
            probs = D ** 2 / np.sum(D ** 2)
            next_center = np.random.choice(N, p=probs)
            centers.append(next_center)
        labels = np.zeros(N, dtype=int)
        for i in range(N):
            labels[i] = np.argmax([K[i, c] for c in centers])
    else:
        raise ValueError("init_method must be 'random' or 'kmeans++'")

    snapshots = []

    for it in range(max_iter):
        cluster_indices = [np.where(labels == c)[0] for c in range(k)]
        intra_K = np.zeros(k)
        for c in range(k):
            idx = cluster_indices[c]
            if len(idx) > 0:
                intra_K[c] = np.sum(K[np.ix_(idx, idx)]) / (len(idx) ** 2)

        new_labels = np.zeros(N, dtype=int)
        for i in range(N):
            best_c = 0
            min_dist = float('inf')
            for c in range(k):
                idx = cluster_indices[c]
                if len(idx) == 0:
                    continue
                term1 = K[i, i]
                term2 = -2 * np.sum(K[i, idx]) / len(idx)
                term3 = intra_K[c]
                dist = term1 + term2 + term3
                if dist < min_dist:
                    min_dist = dist
                    best_c = c
            new_labels[i] = best_c

        snapshots.append(new_labels.copy())
        changed = np.sum(new_labels != labels)
        print(f"Iteration {it + 1}: {changed} points changed.")

        if changed < tol * N:
            print("Converged.")
            break

    labels = new_labels

return labels, snapshots

```

▪ Step 1: Compute intra-cluster kernel terms

For each cluster c :

- Find all indices i assigned to cluster c .
- Compute:

$$\text{intra_K}[c] = \frac{1}{|C_c|^2} \sum_{i,j \in C_c} K[i, j]$$

- This represents the squared norm of the centroid μ_c^ϕ in feature space.

▪ Step 2: Assign new cluster labels

For each data point x_i , compute its distance to each cluster centroid using:

$$\text{dist}(x_i, C_c) = K[i, i] - \frac{2}{|C_c|} \sum_{j \in C_c} K[i, j] + \text{intra_K}[c]$$

- Assign point to the cluster with the **minimum kernel-based distance**.

▪ Step 3: Check convergence

- Count how many points changed labels.
- If `changed < tol * N`, stop early (converged).

```
def kernel_kmeans_pipeline(image_path, n_clusters, gamma_s, gamma_c, init_method='kmeans++'):
    img, h, w = load_image(image_path)
    K = compute_kernel_matrix(img, gamma_s, gamma_c)
    labels, snapshots = kernel_kmeans(K, n_clusters, init_method=init_method)
    file_prefix = f'{os.path.splitext(os.path.basename(image_path))}'
    [0]_kernel_{init_method}_{n_clusters}clusters"
    save_gif_and_png(snapshots, h, w, n_clusters, file_prefix)
```

Executes the full kernel k-means clustering pipeline on an image. Includes feature extraction, kernel computation, clustering, and visualization

spectral clustering

Spectral Clustering is a graph-based clustering algorithm that uses the eigenvectors of the graph Laplacian to embed the data into a lower-dimensional space, where K-Means can then be applied.

```
def construct_laplacian(W, method='normalized'):
    D = np.diag(np.sum(W, axis=1))
    if method == 'ratio':
        return D - W
    elif method == 'normalized':
        D_inv_sqrt = np.diag(1.0 / np.sqrt(np.diag(D)))
        return np.identity(W.shape[0]) - D_inv_sqrt @ W @ D_inv_sqrt
    else:
        raise ValueError("method must be 'normalized' or 'ratio'")
```

- Constructs the **graph Laplacian matrix** used in Spectral Clustering.
- W is the **similarity (adjacency) matrix**, typically derived from a kernel.
- Two common variants: - **Ratio cut (unnormalized)**:

$$L = D - W$$

- Normalized cut (symmetric):

$$L_{\text{sym}} = I - D^{-1/2} W D^{-1/2}$$

Where:

- D : Degree matrix, a diagonal matrix with $D_{ii} = \sum_j W_{ij}$
- I : Identity matrix

```
def spectral_embedding(L, n_clusters, skip_first=True):
    eigvals, eigvecs = eigh(L)
    start_idx = 1 if skip_first else 0
    U = eigvecs[:, start_idx:start_idx + n_clusters]
    return U / np.linalg.norm(U, axis=1, keepdims=True)
```

- Performs **eigen-decomposition** of the Laplacian matrix L to obtain the **embedding matrix U** .
- This embedding captures the geometry of the data graph.
- The first eigenvector is usually skipped (`skip_first=True`) because: - For normalized Laplacian: the first eigenvalue is 0 and corresponds to the trivial constant eigenvector. Output:
- $U \in \mathbb{R}^{N \times k}$ where each row is the projection of a point into the **spectral space**. Why normalize the rows?
- Row-wise normalization ensures that clustering is based on **direction** (angular similarity) in the spectral space.

```
def kmeans(X, k, max_iter=100, tol=1e-4, init_method='kmeans++'):
    N = X.shape[0]
    centers = initialize_kmeans(X, k, init_method)
    labels = np.zeros(N, dtype=int)
    snapshots = []

    for it in range(max_iter):
        dists = np.linalg.norm(X[:, None, :] - centers[None, :, :], axis=2)
        new_labels = np.argmin(dists, axis=1)
        snapshots.append(new_labels.copy())

        if np.sum(new_labels != labels) < tol * N:
            print(f"Converged at iteration {it + 1}")
            break

        labels = new_labels
        for i in range(k):
            if np.any(labels == i):
                centers[i] = X[labels == i].mean(axis=0)

    return labels, snapshots
```

- Applies the **K-Means algorithm** to points in the **spectral embedding space**.
- Used after computing eigenvectors of the Laplacian.
- Initialization supports:
 - `'random'` : randomly choose initial centroids
 - `'kmeans++'` : probabilistic spread-out selection of centroids
- Tracks label changes across iterations for convergence.

```
def spectral_clustering_pipeline(image_path, n_clusters, gamma_s, gamma_c, method='normalized',
init_method='kmeans++'):
    img, h, w = load_image(image_path)
    K = compute_kernel_matrix(img, gamma_s, gamma_c)
    L = construct_laplacian(K, method)
    U = spectral_embedding(L, n_clusters)
    labels, snapshots = kmeans(U, n_clusters, init_method=init_method)
    file_prefix = f"{os.path.splitext(os.path.basename(image_path))}"
    [0]}_spectral_{method}_{init_method}_{n_clusters}clusters"
```

```
save_gif_and_png(snapshots, h, w, n_clusters, file_prefix)
visualize_eigenspace(U, labels, file_prefix)
```

1. Load image and extract its RGB data and shape.
2. Compute the kernel matrix K using both spatial and color similarities.
3. Construct the graph Laplacian (L) using ratio or normalized method.
4. Compute eigenvectors of the Laplacian and extract the top $n_{clusters}$ ones as embedding matrix U .
5. Apply K-Means in the eigenspace to cluster pixels.
6. Save results as a .gif (animated clustering) and .png (final output).
7. Visualize the eigenspace (2D or 3D scatter plot of eigenvectors with cluster colors)

Part4: Experiments on the coordinates in the eigenspace

visualize the coordinates of each data point (e.g., image pixel) in that eigenspace, colored by its final cluster assignment.

```
def visualize_eigenspace(U, labels, file_prefix, output_dir='./eigenspace_vis'):
    os.makedirs(output_dir, exist_ok=True)
    k = U.shape[1]
    fig_path = os.path.join(output_dir, f"{file_prefix}_eigenspace.png")
    if k == 2:
        plt.figure(figsize=(8, 6))
        for c in np.unique(labels):
            plt.scatter(U[labels == c, 0], U[labels == c, 1], label=f"Cluster {c}", s=5)
        plt.xlabel("Eigenvector 1")
        plt.ylabel("Eigenvector 2")
        plt.title("Eigenspace Visualization (2D)")
        plt.legend()
        plt.grid(True)
        plt.savefig(fig_path)
        plt.close()
        print(f"Eigenspace plot saved to: {fig_path}")
    elif k == 3:
        fig = plt.figure(figsize=(10, 8))
        ax = fig.add_subplot(111, projection='3d')
        for c in np.unique(labels):
            ax.scatter(U[labels == c, 0], U[labels == c, 1], U[labels == c, 2], label=f"Cluster {c}",
                      s=5)
        ax.set_title("Eigenspace Visualization (3D)")
        ax.set_xlabel("Eigenvector 1")
        ax.set_ylabel("Eigenvector 2")
        ax.set_zlabel("Eigenvector 3")
        ax.legend()
        plt.savefig(fig_path)
        plt.close()
        print(f"Eigenspace 3D plot saved to: {fig_path}")
    else:
        print(f"[Warning] Eigenspace visualization only supported for k=2 or 3 (got k={k}). Skipped.")
```

- `U` : Spectral embedding matrix (shape: $N \times k$), where each row is the coordinate of a data point in eigenspace.
- `labels` : Cluster assignments for each point.
- `file_prefix` : Output file naming base.
- `output_dir` : Directory to save the image.

Output

- Saves a `.png` image showing:
 - **2D scatter plot** (if $k = 2$)
 - **3D scatter plot** (if $k = 3$)
- Each point is colored by its cluster label.

```
def spectral_clustering_pipeline(...):
    ...
    save_gif_and_png(snapshots, h, w, n_clusters, file_prefix)
    visualize_eigenspace(U, labels, file_prefix)
```

```
def labels_to_rgb(labels, n_clusters):
    colormap = plt.get_cmap("tab10", n_clusters)
    colors = (colormap(np.arange(n_clusters))[:, :3] * 255).astype(np.uint8)
    return colors[labels]
```

```
def save_gif_and_png(snapshots, h, w, n_clusters, file_prefix, output_dir="../output"):
    os.makedirs(output_dir, exist_ok=True)
    gif_frames = []
    for labels in snapshots:
        rgb_img = labels_to_rgb(labels, n_clusters).reshape(h, w, 3)
        gif_frames.append(rgb_img)

    gif_path = os.path.join(output_dir, f"{file_prefix}.gif")
    imageio.mimsave(gif_path, gif_frames, format='GIF', duration=0.5)
    print(f"Saved GIF to {gif_path}")

    png_path = os.path.join(output_dir, f"{file_prefix}_final.png")
    final_frame = gif_frames[-1]
    Image.fromarray(final_frame).save(png_path)
    print(f"Saved final PNG to {png_path}")
```

- Converts each label map (from `snapshots`) to an RGB image.
- Creates an animated `.gif` of the entire clustering process.
- Saves the final frame (last iteration) as a `.png`.

```
def run_all_configs(image_path, gamma_s, gamma_c, n_clusters_list, init_methods, modes,
                    laplacian_methods):
    for mode in modes:
        for init_method in init_methods:
            for n_clusters in n_clusters_list:
                if mode == 'kernel':
                    print(f"\n[Kernel K-Means] clusters={n_clusters}, init={init_method}")
                    try:
                        kernel_kmeans_pipeline(
                            image_path=image_path,
                            n_clusters=n_clusters,
```

```

        gamma_s=gamma_s,
        gamma_c=gamma_c,
        init_method=init_method
    )
except Exception as e:
    print(f"Error in kernel_kmeans_pipeline: {e}")

elif mode == 'spectral':
    for lap_method in laplacian_methods:
        print(f"\n[Spectral Clustering] clusters={n_clusters}, init={init_method},"
laplacian={lap_method}")
        try:
            spectral_clustering_pipeline(
                image_path=image_path,
                n_clusters=n_clusters,
                gamma_s=gamma_s,
                gamma_c=gamma_c,
                method=lap_method,
                init_method=init_method
            )
        except Exception as e:
            print(f"Error in spectral_clustering_pipeline: {e}")

else:
    print(f"Unknown mode: {mode}")

```

Runs **batch experiments** for all combinations of clustering mode, number of clusters, initialization methods, and Laplacian types.

Parameters:

- `image_path` : Path to the input image.
- `gamma_s` , `gamma_c` : Kernel parameters for spatial and color similarity.
- `n_clusters_list` : List of values for number of clusters (e.g., `[2, 3, 4]`).
- `init_methods` : List of initialization strategies (e.g., `['random', 'kmeans++']`).
- `modes` : `['kernel', 'spectral']` for selecting clustering algorithms.
- `laplacian_methods` : Applicable only if mode is `'spectral'` (e.g., `['ratio', 'normalized']`). Features:
 - Each configuration will:
 - Print progress and configuration info.
 - Attempt to run the full pipeline and handle exceptions gracefully.
 - Output includes `.gif` , `.png` , and optional eigenspace plots.

2. Experiments settings and results(30%) & discussion (20%)

Part1+Part2 + Part3

TA Can See the GIF images version in this link

```
if __name__ == "__main__":
    gamma_s = 0.001
    gamma_c = 0.001
    n_clusters_list = [2, 3, 4] # Part2: Try more clusters (5%)
    init_methods = ['random', 'kmeans++'] # Part3: Try different initializations. (5%)
    modes = ['kernel', 'spectral']
    laplacian_methods = ['ratio', 'normalized']

    for i in range(1, 3):
        image_path = f"./image{i}.png"
        run_all_configs(
            image_path=image_path,
            gamma_s=gamma_s,
            gamma_c=gamma_c,
            n_clusters_list=n_clusters_list,
            init_methods=init_methods,
            modes=modes,
            laplacian_methods=laplacian_methods
        )
```

Image1

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	random	random	random
2 clusters			
3 clusters			
4 clusters			

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	k-means++	k-means++	k-means++

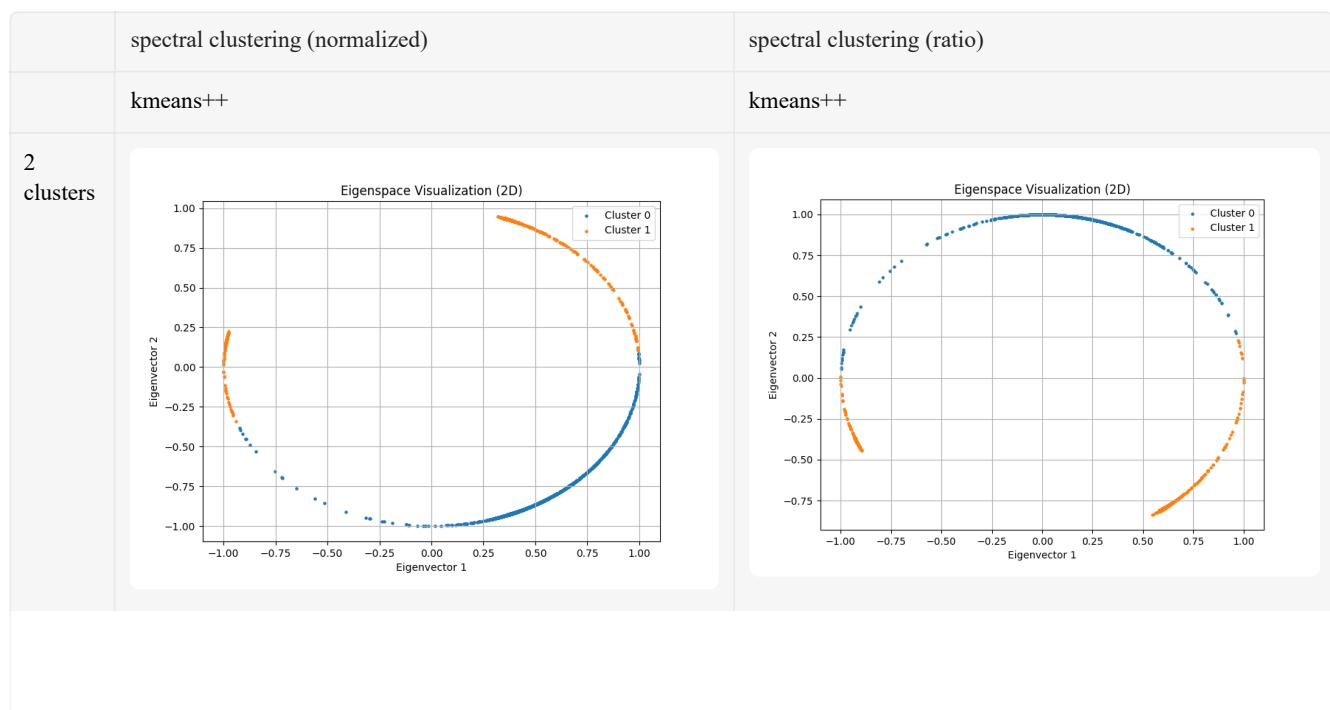
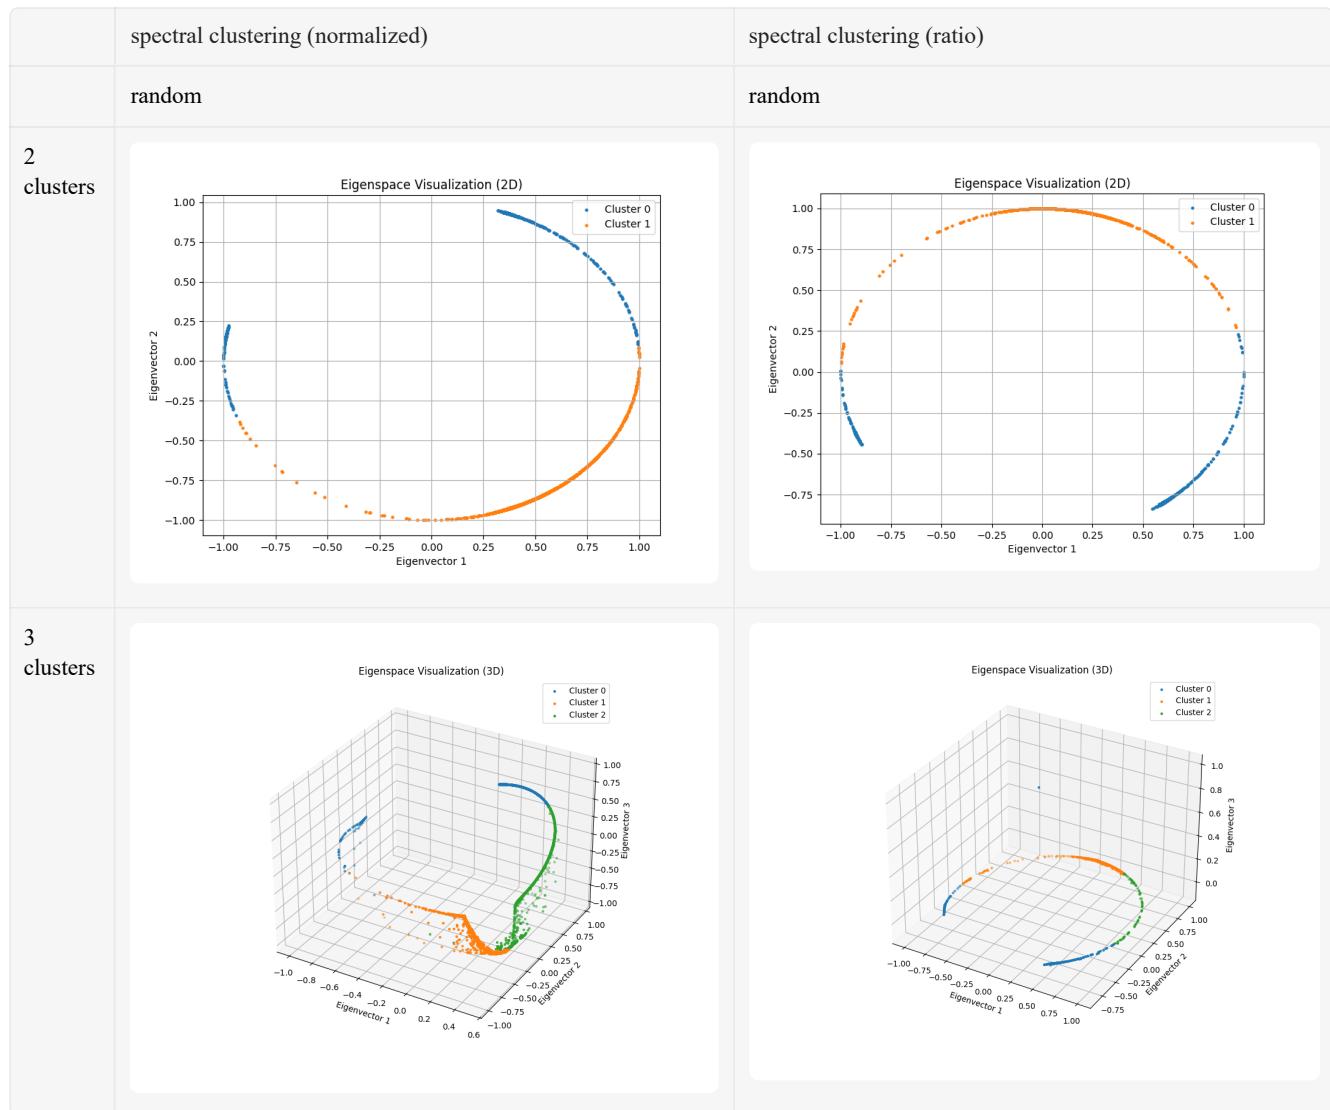
	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
2 clusters			
3 clusters			
4 clusters			

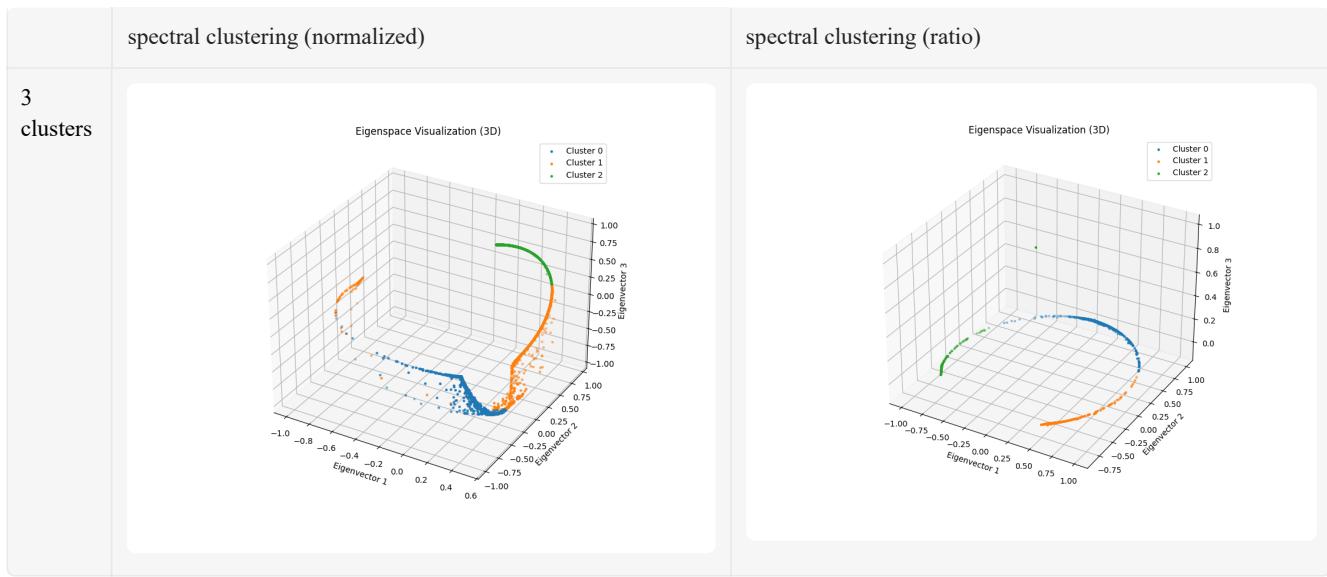
	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	random	random	random
2 clusters			
3 clusters			
4 clusters			

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	k-means++	k-means++	k-means++
2 clusters			
3 clusters			
4 clusters			

Part4

To observe how well-separated the clusters are in the spectral embedding space (formed by the Laplacian's eigenvectors). Better separability in eigenspace often leads to more reliable clustering.

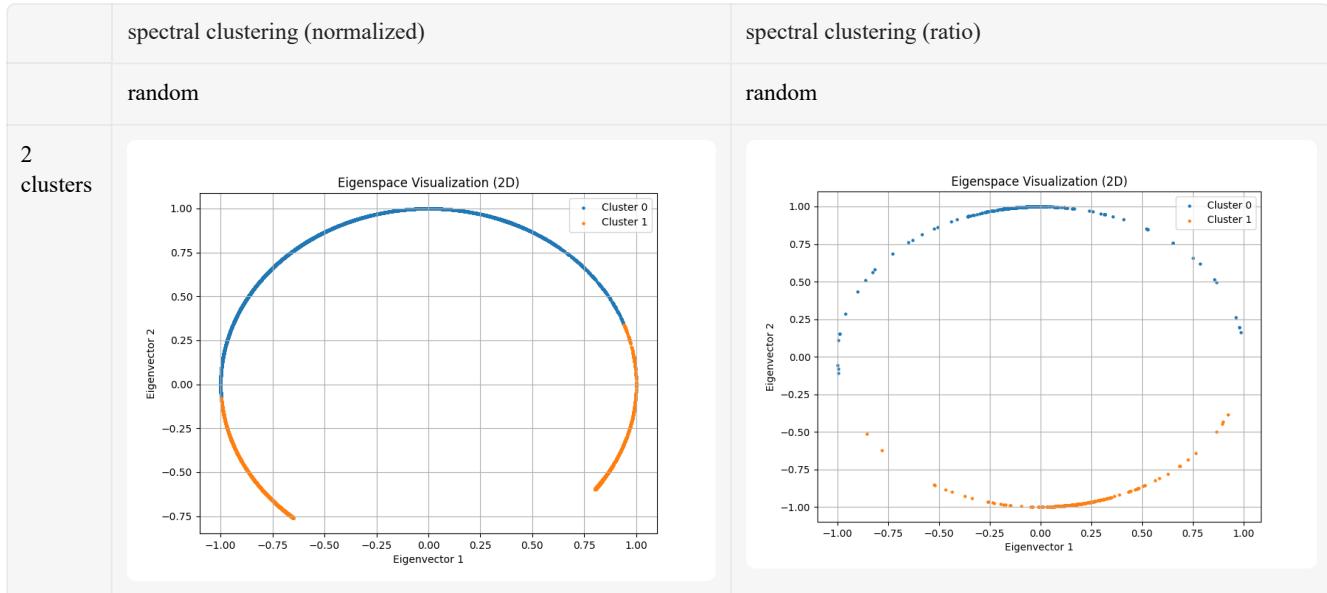


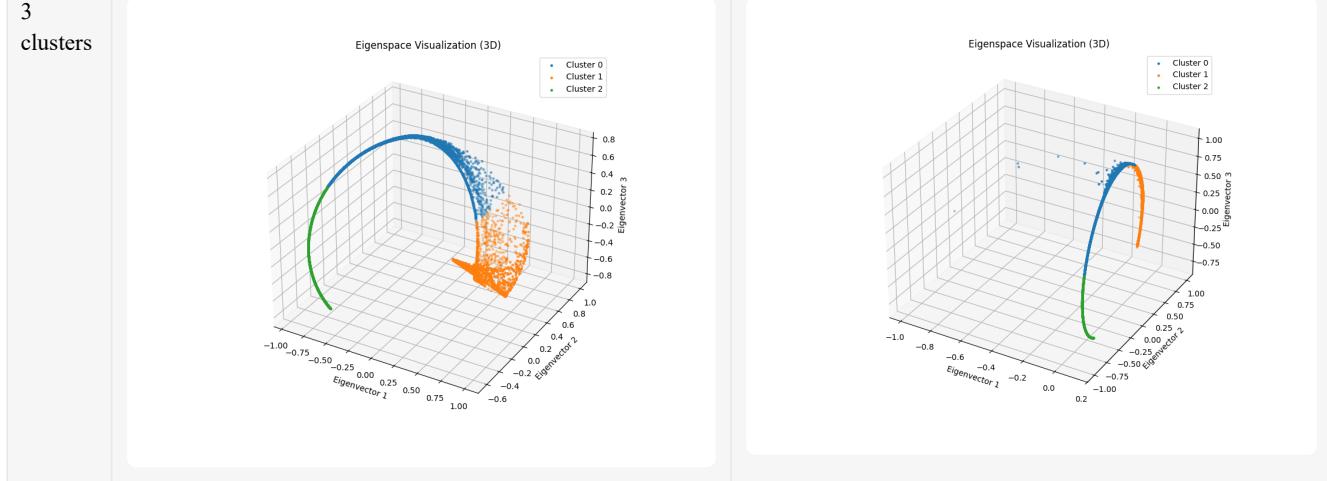
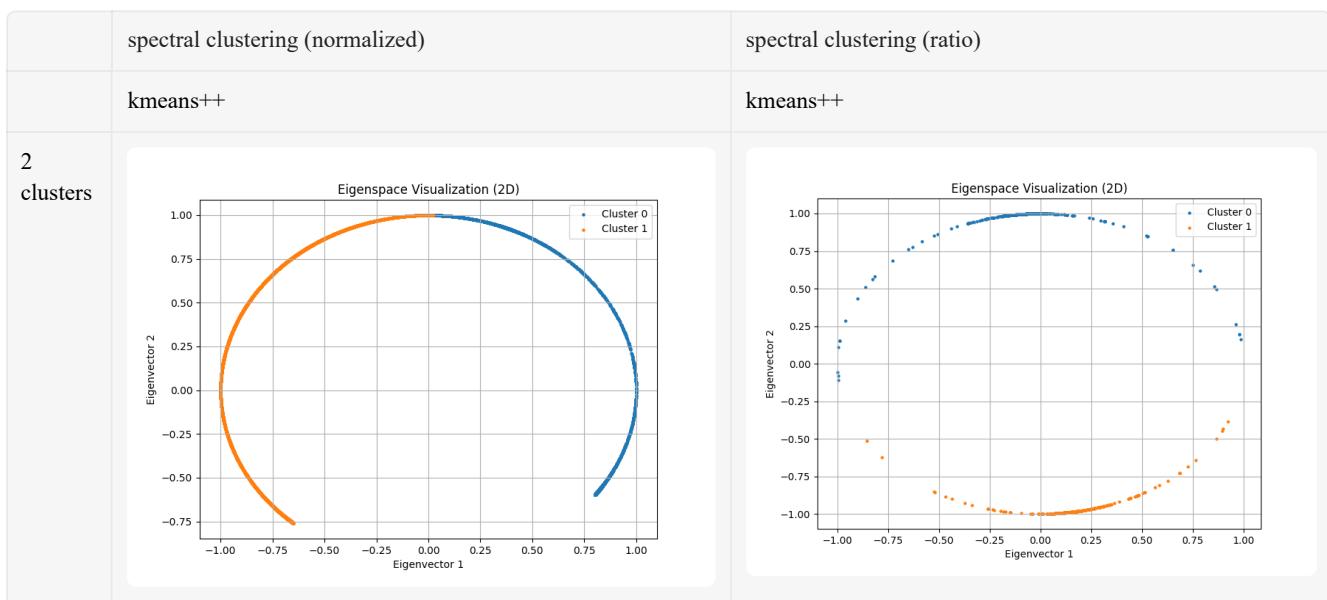
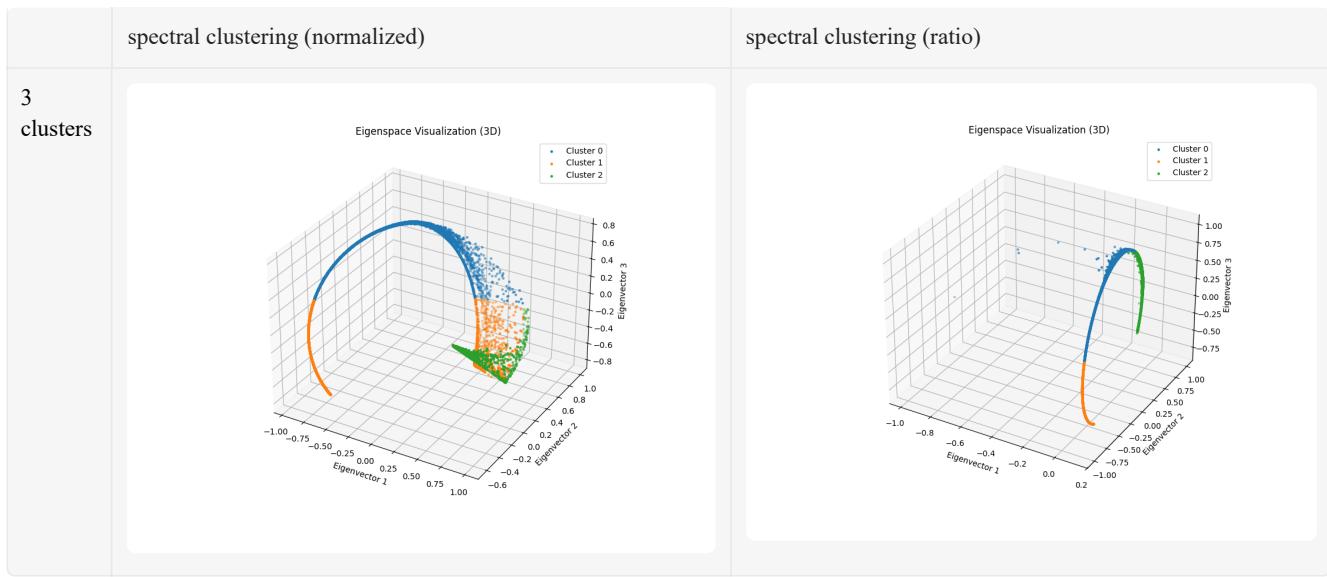


Aspect	Spectral (Normalized)	Spectral (Ratio)
2 Clusters	Strong structure, circular arc	Same; slightly tighter clusters
3 Clusters	Smooth curved embedding, good	More compact but slightly sparse
k-means++ vs Random	Improves clarity	Greatly improves density & balance

- The **eigenspace structure of Image1** is **strongly curved**, which makes spectral clustering ideal.
- Normalized Laplacian** provides smoother embeddings.
- Ratio Laplacian** benefits more from **k-means++**, especially at higher cluster counts.

image2





Setting	Normalized Laplacian	Ratio Laplacian
2 Clusters	Arc-shaped, clean boundary	Similar arc, more sensitive to init
3 Clusters	Hook-shaped 3D arc, moderate overlap	Good shape but less balanced w/ random
k-means++ Effect	Refines boundaries, improves clarity	Dramatically improves compactness & separation

- Image2's spectral embeddings **form clear low-dimensional manifolds**.
- **Spectral Clustering (Ratio)** can match or exceed normalized performance when paired with **k-means++**.
- Choosing proper initialization is crucial when embeddings are **sparse or uneven**.

3. Observations and discussion (20%)

Compare the performance between different clustering methods. (8%)

image1

2 Clusters

Random Initialization

- **Kernel K-Means:** Clusters appear noisy, some irregularities near boundaries.
- **Spectral Clustering (Normalized/Ratio):** Cleaner segmentation; objects are more clearly separated, especially the background and main object. **K-Means++ Initialization**
- Improved consistency across all methods.
- **Spectral Clustering (Normalized)** maintains the most visually stable and coherent segmentation.
- **Kernel K-Means** becomes noticeably better, with reduced noise and clearer regions.

Spectral Clustering (especially normalized) effectively captures global structures even with just two clusters. **k-means++** helps eliminate instability from random starts.

3 Clusters

Random Initialization

- **Kernel K-Means:** More fragmentation observed. One region appears over-segmented.
- **Spectral Clustering (Normalized):** Provides distinguishable separation of three regions.
- **Spectral Clustering (Ratio):** Similar to normalized, though slightly less defined.

K-Means++ Initialization

- Clear performance improvement in all methods.
- **Spectral Clustering (Normalized)** shows the sharpest boundaries and consistent cluster shapes.
- **Kernel K-Means** is more balanced than in the random case.

With more clusters, initialization plays a bigger role. Spectral Clustering retains robustness and interpretability.

4 Clusters

Random Initialization

- All methods show some fragmentation or over-segmentation.
- **Spectral (Normalized)** still outperforms the rest, though boundary clarity is slightly reduced.
- **Kernel K-Means** has the most scattered small regions.

K-Means++ Initialization

- **Spectral Clustering (Normalized)** shows a clear and meaningful 4-region split.
- **Spectral (Ratio)** is still acceptable but slightly noisier.
- **Kernel K-Means** improves compared to random, but cluster shapes are less smooth.

When the number of clusters increases, Spectral Clustering with normalization remains the most reliable. Kernel K-Means becomes less stable and more sensitive to initialization.

image2

2 Clusters

Random Initialization

- **Kernel K-Means:**
 - Clean split with large uniform regions.
 - Some subtle noise on boundaries.
- **Spectral Clustering (Normalized & Ratio):**
 - More sensitive to edge textures and finer details.
 - Resulting segmentation includes scattered fragments.
 - Ratio method appears slightly smoother than normalized.

K-Means++ Initialization

- **Kernel K-Means:**
 - Still clean and consistent; slightly more balanced than with random.
- **Spectral Clustering:**
 - Still noisy in fine-textured regions, but more stable than random.
 - Both normalized and ratio produce very similar results.

Kernel K-Means is more spatially coherent under 2-cluster scenarios, while Spectral Clustering is more sensitive to local texture details. Initialization method does not heavily affect outcome at 2 clusters, but `k-means++` offers slightly cleaner results.

3 Clusters

Random Initialization

- **Kernel K-Means:**
 - Good separation of major areas, but some mixed boundaries.
- **Spectral Clustering:**
 - Segment boundaries align more with texture transitions.
 - Noticeable patchiness, especially in normalized version.

K-Means++ Initialization

- **Kernel K-Means:**
 - Reduced patch noise, better-defined cluster shapes.
- **Spectral Clustering:**
 - Less chaotic than random; smoother transitions.
 - Normalized and ratio still differ slightly in region compactness.

With more clusters, the effect of initialization becomes more noticeable. Spectral methods capture more nuanced structures but are prone to over-segmentation without proper tuning.

4 Clusters

Random Initialization

- **Kernel K-Means:**
 - Cluster regions look natural, but smaller details may be under-represented.
- **Spectral Clustering:**

- Both normalized and ratio highlight very fine-grained structures.
- More scattered pixel clusters, which might be undesirable for object-level segmentation.

K-Means++ Initialization

- **Kernel K-Means:**
 - Significantly improved — clearer separation and compact clusters.
- **Spectral Clustering:**
 - Ratio version becomes more structured and readable.
 - Normalized still prone to noise in textured regions, but better than random.

With higher cluster counts, `k-means++` is essential to stabilize the results. Spectral Clustering (Ratio) becomes more visually coherent than Normalized in high-cluster settings.

Compare the execution time of different settings. (8%)

image1

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	random	random	random
2 clusters	40.37 sec	312.72 sec	173.45 sec
3 clusters	30.43 sec	176.54 sec	200.60 sec
4 clusters	84.31 sec	209.88 sec	197.34 sec

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	<code>k-means++</code>	<code>k-means++</code>	<code>k-means++</code>
2 clusters	45.47 sec	208.94 sec	150.79 sec
3 clusters	41.60 sec	210.32 sec	173.40 sec
4 clusters	79.79 sec	218.64 sec	152.71 sec

Observations

1. General Trends

- **Kernel K-Means** is significantly faster across all settings.
- **Spectral Clustering (Normalized)** is the slowest method, particularly for `2 clusters` and `k-means++`.
- **Spectral Clustering (Ratio)** lies between the two in terms of runtime, generally faster than normalized but slower than kernel k-means.

2. Effect of Initialization Method

- `k-means++ initialization` increases runtime slightly for **Kernel K-Means**, likely due to the additional computation of better-initialized centers.
- Surprisingly, for **Spectral Clustering**, `k-means++` slightly **reduces or stabilizes runtime** compared to random — likely because better initialization reduces k-means iterations after spectral embedding.

3. Effect of Cluster Number

- For all methods, increasing the number of clusters does **not always linearly increase runtime**.

- **Kernel K-Means:** runtime fluctuates, e.g., 3 clusters takes **less time** than 2 clusters in the random setting — likely due to early convergence.
 - **Spectral Clustering (Normalized):** runtime increases with more clusters, especially in k-means++ .
 - **Spectral Clustering (Ratio):** similar pattern but with smaller increases.
-

- **Kernel K-Means** is clearly the most efficient method, making it suitable for time-sensitive applications or larger images.
- **Spectral Clustering** methods, while slower, may offer better segmentation quality — a trade-off between accuracy and speed.
- Choosing k-means++ may slightly increase cost for Kernel K-Means but improves convergence in Spectral Clustering.
- **Spectral Clustering (Ratio)** provides a good middle ground in terms of both **runtime** and **segmentation smoothness**.

image2

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	random	random	random
2 clusters	19.10 sec	203.68 sec	176.47 sec
3 clusters	19.68 sec	206.87 sec	153.53 sec
4 clusters	20.88 sec	207.15 sec	164.27

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	k-means++	k-means++	k-means++
2 clusters	15.05 sec	209.93sec	166.47 sec
3 clusters	44.16 sec	196.51 sec	167.22 sec
4 clusters	36.29 sec	196.04 sec	170.13 sec

Observations

1. Overall Performance

- **Kernel K-Means** is significantly faster than both spectral clustering variants in all configurations.
 - **Spectral Clustering (Normalized)** has the highest runtime consistently (~196–210 seconds).
 - **Spectral Clustering (Ratio)** performs slightly faster than normalized in most settings, especially under random .
-

2. Effect of Initialization Method

Kernel K-Means:

- **Unusual behavior:**
 - k-means++ is **faster** than random in 2 clusters , but **slower** for 3 and 4 clusters .
 - This may be due to variability in early convergence behavior depending on initialization.

Spectral Clustering:

- The choice between random and k-means++ **has minimal impact** on execution time.
 - This is expected, since most of the time is spent computing the Laplacian and its eigenvectors, not in k-means.

3. Effect of Cluster Number

Kernel K-Means:

- Runtime increases slightly with more clusters under `random`.
- Under `k-means++`, runtime spikes for `3 clusters` (44.16 sec) and drops again for `4 clusters` — possibly due to convergence complexity at `k=3`.

Spectral Clustering:

- Both **Normalized** and **Ratio** variants maintain **stable runtime** across different cluster numbers.
- This shows that the **eigen-decomposition dominates** runtime and is not strongly affected by `k`.

- **Kernel K-Means** is extremely efficient on Image2, completing most runs under **45 seconds**.
- **Spectral Clustering** is significantly slower, taking **~3 to 10× longer**, regardless of initialization or cluster count.
- **Spectral (Ratio)** remains a viable trade-off between runtime and structure-aware segmentation.
- `k-means++` helps Kernel K-Means only in specific settings; for spectral methods, it has **negligible effect on time**.

Anything you want to discuss. (4%)

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	random	random	random
2 clusters	15	4	4
3 clusters	11	9	23
4 clusters	38	12	19

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	<code>k-means++</code>	<code>k-means++</code>	<code>k-means++</code>
2 clusters	22	4	4
3 clusters	14	12	2
4 clusters	33	7	4

- **Spectral Clustering** (both normalized and ratio) consistently converges **much faster** (mostly within 2–12 iterations) than **Kernel K-Means**, which often requires **10–40+ iterations**.
- **Normalized Spectral Clustering** is especially stable and predictable, showing **almost constant convergence** around 4–12 iterations regardless of `k`.
- **Kernel K-Means** is **more sensitive** to:
 - the number of clusters (`k`)
 - the initialization method (`k-means++` generally helps but not always)

image2

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	random	random	random
2 clusters	11	8	3
3 clusters	13	13	9
4 clusters	17	20	62

	kernel k-means method	spectral clustering (normalized)	spectral clustering (ratio)
	k-means++	k-means++	k-means++
2 clusters	8	21	2
3 clusters	41	9	8
4 clusters	34	10	10

- **Spectral Clustering (Ratio)** shows **extreme sensitivity** to cluster number and initialization:
 - It converges **very quickly** in some cases (e.g., 2 clusters, `k-means++` : 2 iterations),
 - But can take **very long** in others (e.g., 4 clusters, `random` : 62 iterations).
- **Spectral Clustering (Normalized)** is generally more **stable** in convergence speed, mostly staying below 20 iterations.
- **Kernel K-Means** remains **variable**:
 - `k-means++` sometimes reduces iterations (e.g., 2 clusters: from 11 → 8),
 - But may also **increase** iteration count unexpectedly (e.g., 3 clusters: 13 → 41).